

# R Machine Learning Manual

*Abu Nayeem*

*September 23, 2014*

**Introduction** This Manual is meant to consolidate my knowledge in machine learning using the R package. The first two steps of machine learning is the following:

- What is the predictor variable [factor vs. numeric]
- How do you clean the dataset without removing valuable information?

Of course the standard practice of machine learning involves creating a training test [create model], cross validation set [test model], and a testing set [final unadulterated data]. This sort of testing procedure assures that the model does not face the issue of overfitting. This example will be a classification problem, but a numeric learning problem follow similar procedure.

**Executive Summary** Technology has focused on developing health tools and gadgets to record how much training a person has done in a specific period of time. However, almost no research has been done in developing tools or models to give the trainer feedback on how well he has been performing exercises. This project is oriented in calculating a machine learning algorithm to determine whether a weight lifting trainer performed the exercise well or made an error in the execution.

**DataSet** The data set used for the model comes from the [Groupware@LES](http://groupware.les.inf.puc-rio.br/har) from their Human Activity Recognition project. They performed a study to analyze how well a Weight Lifting Exercise was executed. Each trainer was given a sensor for his glove, belt, dumbbell and arm-band. These are tools used by every weight lifting trainer so the original exercises maintain integrity.

Each trainer was asked to perform weight lifting in a particular manner. First, to do it perfectly as ideally described. Second, throwing the elbows to the front. Third, lifting the dumbbell half way. Fourth, lowering the dumbbell halfway. Finally, throwing the hips to the front. In each exercise performed, the sensors recorded the movements and rotations, including max accelerations, min accelerations, averages, kurtosis, between others.

You can learn more here <http://groupware.les.inf.puc-rio.br/har>

## Preprocessing

```
set.seed(234)
library(caret) # the power horse function; loads ggplot2 automatically
library(doMC) # enable parallel computing; loads parallel & iterators
library(nnet) # for neural networking and multi-nomial log regression models
library(randomForest) # random forest strategy
library(kernlab) # allows plenty of tools of dimension reduction and such
library(e1071) # allows more features but is needed for boosting models
library(plyr) # data table operations
library(dplyr) # data operations plus
library(gbm) # general boosting method
library(corrplot) # fancy correlation plot
```

```
library(AppliedPredictiveModeling)
library(foreach) # used in random forest algorithm
library(doParallel) # Parallel Processing
library(ipred) # needed for treebagging
library(rpart) # for rpart but it failed in this example
registerDoMC(cores = 2) # register the number of cores to parallel process
date() #set date
```

## Loading libraries

```
## [1] "Fri Sep 26 00:13:16 2014"
```

```
# Selecting the definition of NA string was defined via post-analysis
trainingfile <- 'http://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv'
training <- read.csv(trainingfile, na.strings = c("NA", "#DIV/0!"))
training <- tbl_df(training) # this data table is smoother
testingfile <- 'http://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv'
testing <- read.csv(testingfile, na.strings = c("NA", "#DIV/0!"))
testing <- tbl_df(testing)
```

## Extraction

**Setup** An unresolved debate is do you keep the cross validation pristine prior to implementing data cleaning techniques. I recommend testing your algorithm in both options to see if it actually makes a difference. The cross validation process lose precious data which can be used to create a prediction model. Let's give it the worse case scenario where cross-validation set is also treated as pristine. Note: you cannot fundamentally change the training set because the testing set is still raw , so you should take that to account.

Splitting training set into a smaller training set and cross-validation set

```
inTrain <- createDataPartition(y = training$classe, p = 0.8, list = FALSE)
smalltraining <- training[inTrain, ]
crossvalidation <- training[-inTrain, ]
```

## Data Cleaning


A) **Basics-** Make basic assessment that need to be done

```
dim(training)
str(training)
summary(training)
```

B) **Handling missing values:** note you should set the strings for missing upon retrieving the data

Plotting of missing values:

```
qplot(1, colSums(is.na(smalltraining))/dim(smalltraining)[1],
      geom = 'jitter',
      main = '% of missing values per variable',
      xlab = '', ylab = '% missing values') # visualization of missing values
```



Accumulative method Removal:

```
colSums(is.na(smalltraining)) # now we see the number of missing values in columns
                              # and see if they are significant for removal
NonNAIndex <- which(colSums(is.na(smalltraining)) > 0)
# this extracts the column index missing variable with number
RemoveNA <- smalltraining[, -NonNAIndex]
# Create new data frame that remove columns that had missing values
```

Threshold Method: you choose the tolerable percentage and if above, remove the columns.

```
NonNAIndex <- which(apply(smalltraining, 2,
                          function(x) {sum(is.na(x))}) > 0.5 * dim(smalltraining)[1])
RemoveNA <- smalltraining[, -NonNAIndex]

# alternative to above but more function like
NA_threshold <- 0.50
nTrain <- nrow(smalltraining)
i <- 1
while(i < ncol(smalltraining)) {
  nNA <- sum(is.na(smalltraining[,i]))
  if((nNA/nTrain) >= NA_threshold) {
    NonNAIndex <- c(removeCols, i)
  }
  i <- i + 1
}
RemoveNA <- smalltraining[, -NonNAIndex]
```

### C) Removing Uninteresting Features:

Removing Unrelated Features:

```
# choose the columns that may be useful for analysis  
compacttraining <- select(RemoveNA, 2, 8:60)
```

Removing Near Zero variance features:

```
# this checks if all columns have close to zero variance  
# the saveMetric provide heuristic information of each column which is REALLY useful  
Nzv <- nearZeroVar(compacttraining,saveMetrics=TRUE)  
Nzv # all false, so no columns will be removed
```

##	freqRatio	percentUnique	zeroVar	nzv
## user_name	1.117	0.03822	FALSE	FALSE
## roll_belt	1.103	7.54188	FALSE	FALSE
## pitch_belt	1.032	11.14084	FALSE	FALSE
## yaw_belt	1.017	11.77782	FALSE	FALSE
## total_accel_belt	1.052	0.17199	FALSE	FALSE
## gyros_belt_x	1.038	0.83445	FALSE	FALSE
## gyros_belt_y	1.119	0.42678	FALSE	FALSE
## gyros_belt_z	1.062	1.04465	FALSE	FALSE
## accel_belt_x	1.005	1.02554	FALSE	FALSE
## accel_belt_y	1.105	0.87267	FALSE	FALSE
## accel_belt_z	1.074	1.82814	FALSE	FALSE
## magnet_belt_x	1.028	1.98102	FALSE	FALSE
## magnet_belt_y	1.085	1.85362	FALSE	FALSE
## magnet_belt_z	1.016	2.84095	FALSE	FALSE
## roll_arm	55.204	16.04561	FALSE	FALSE
## pitch_arm	84.562	18.50436	FALSE	FALSE
## yaw_arm	34.679	17.40238	FALSE	FALSE
## total_accel_arm	1.050	0.42041	FALSE	FALSE
## gyros_arm_x	1.012	4.05121	FALSE	FALSE
## gyros_arm_y	1.415	2.36958	FALSE	FALSE
## gyros_arm_z	1.108	1.51602	FALSE	FALSE
## accel_arm_x	1.015	4.87292	FALSE	FALSE
## accel_arm_y	1.131	3.38238	FALSE	FALSE
## accel_arm_z	1.120	4.94936	FALSE	FALSE
## magnet_arm_x	1.123	8.47825	FALSE	FALSE
## magnet_arm_y	1.039	5.48443	FALSE	FALSE
## magnet_arm_z	1.046	8.00688	FALSE	FALSE
## roll_dumbbell	1.196	85.47678	FALSE	FALSE
## pitch_dumbbell	2.515	83.41296	FALSE	FALSE
## yaw_dumbbell	1.000	84.98630	FALSE	FALSE
## total_accel_dumbbell	1.075	0.26116	FALSE	FALSE
## gyros_dumbbell_x	1.023	1.51602	FALSE	FALSE
## gyros_dumbbell_y	1.246	1.71985	FALSE	FALSE
## gyros_dumbbell_z	1.061	1.27397	FALSE	FALSE
## accel_dumbbell_x	1.046	2.61800	FALSE	FALSE
## accel_dumbbell_y	1.145	2.89827	FALSE	FALSE
## accel_dumbbell_z	1.189	2.57341	FALSE	FALSE
## magnet_dumbbell_x	1.080	6.92401	FALSE	FALSE

## magnet_dumbbell_y	1.157	5.27422	FALSE	FALSE
## magnet_dumbbell_z	1.032	4.27416	FALSE	FALSE
## roll_forearm	12.008	12.73330	FALSE	FALSE
## pitch_forearm	63.120	17.47245	FALSE	FALSE
## yaw_forearm	14.748	11.77782	FALSE	FALSE
## total_accel_forearm	1.117	0.42678	FALSE	FALSE
## gyros_forearm_x	1.017	1.80903	FALSE	FALSE
## gyros_forearm_y	1.033	4.59902	FALSE	FALSE
## gyros_forearm_z	1.101	1.87910	FALSE	FALSE
## accel_forearm_x	1.013	5.00669	FALSE	FALSE
## accel_forearm_y	1.132	6.25518	FALSE	FALSE
## accel_forearm_z	1.008	3.60533	FALSE	FALSE
## magnet_forearm_x	1.015	9.41461	FALSE	FALSE
## magnet_forearm_y	1.203	11.79056	FALSE	FALSE
## magnet_forearm_z	1.044	10.42105	FALSE	FALSE
## classe	1.469	0.03185	FALSE	FALSE

```
Nzv <- nearZeroVar(compacttraining,saveMetrics=FALSE)
```


```
# if there was columns to be removed this would be used
compacttraining <- compacttraining[ ,-Nzv]
```

D) **Removing Highly Correlated Variables:** For numerical/integer columns only

Plotting Correlated Variables:

```
corData<- cor(compacttraining[ ,c(2:53)])
corrplot(corData,
          title = "Corr,per eigenVectors",
          order = "AOE",
          method = "color",
          type = "lower",
          tl.cex = 0.6 ) # plot to have a look to correlations
```

## Corr,per eigenvectors



Removing Correlated Variables (Manual Method)

```
M <- abs(cor(compacttraining[,c(2:53)])) # create a correlation matrix
diag(M) <- 0 # by default the diagonals are one so we make them equal zero
which(M > 0.8, arr.ind = TRUE) # displays correlated pairs names
which(M > 0.8, arr.ind = FALSE) # displays the column numbers of each match
# you can remove certain pairs manually
descriptivetraining <- select(compacttraining,
  -c(magnet_arm_y, pitch_dumbbell, yaw_dumbbell, accel_arm_x,
    gyros_arm_y, pitch_belt, accel_belt_x, yaw_belt, total_accel_belt,
    accel_belt_y, accel_belt_z, gyros_forearm_y,
    gyros_dumbbell_z, gyros_dumbbell_x)) #40 variables left
```

Alternative Removal Method: Note this method had 42 variables left

```
descrCor <- cor(compacttraining[,c(2:53)])
highlyCorDescr <- findCorrelation(descrCor, cutoff = 0.8)
descriptivetraining <- compacttraining[, -highlyCorDescr] #42 variables left
```

Removing high reasonable skewness [Not useful in predictions]: remember numerical columns only

```
factordescriptivetrain <- descriptivetraining[, c(1,40)] # separate non numerical variable
numdescriptivetrain <- descriptivetraining[, -c(1,40)] # separate numerical variables
NonskewIndex <- which(apply(numdescriptivetrain, 2,
```

```

function(x) abs(skewness(x)) > 6)) # find skewed volumms
numdescriptivetrain <- numdescriptivetrain[, -NonskewIndex] # remove skewed columns
cleandata <- cbind(factordescriptivetrain,numdescriptivetrain) # Combine to create clean data

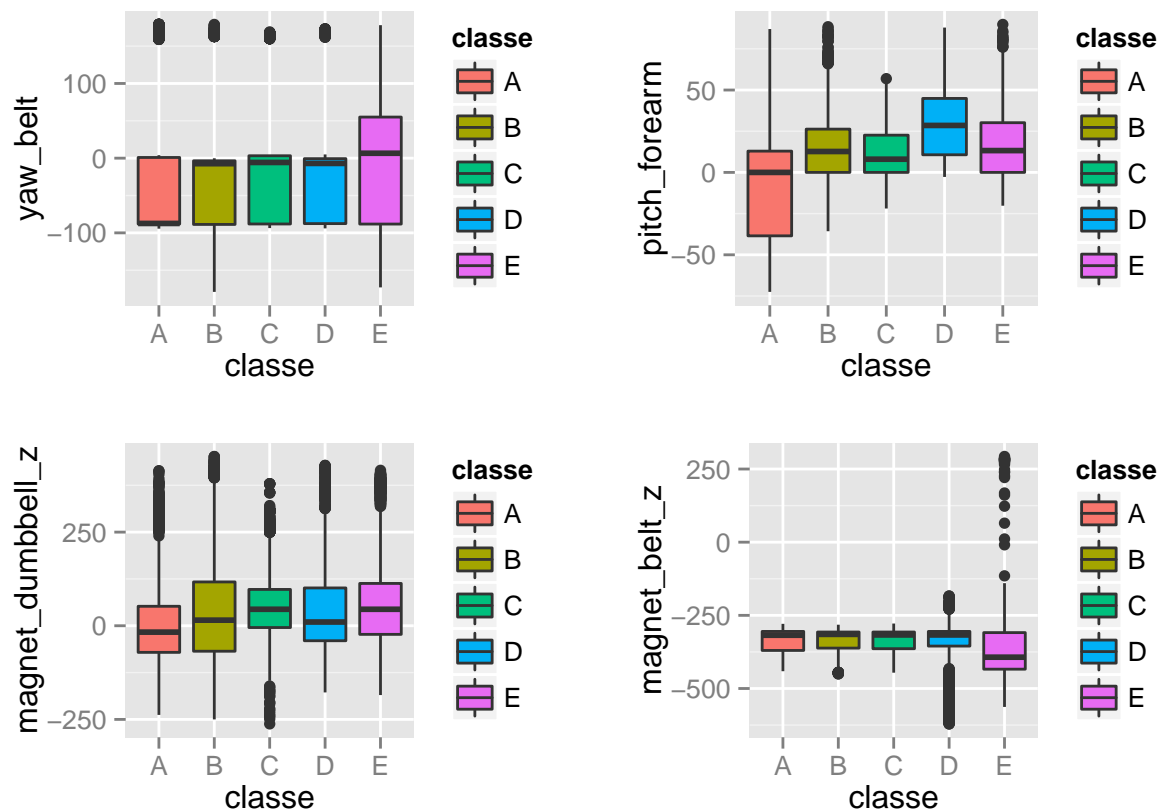
```

E) **Exploratory Analysis:** [if you have an idea which variables are of concern]

```

require(gridExtra)
require(ggplot2)
p1 <- qplot(classe,yaw_belt,geom="boxplot",data=smalltraining,fill=classe)
p2 <- qplot(classe,pitch_forearm,geom="boxplot",data=smalltraining,fill=classe)
p3 <- qplot(classe,magnet_dumbbell_z,geom="boxplot",data=smalltraining,fill=classe)
p4 <- qplot(classe,magnet_belt_z,geom="boxplot",data=smalltraining,fill=classe)
grid.arrange(p1,p2,p3,p4, ncol=2)

```



```

colIndex <- colnames(cleandata) #38 variables should be remaining
check<-smalltraining[,colIndex]; check
# the colnames should be identical to that of cleandata

```

**Complete the Column Index**

**Training Models Pre-training:** Loading- You want to save your results so you don't need to constantly repeat analysis. Also note the rpart method failed.

```
if(file.exists("Machine Learning.RData")) {
  load("Machine Learning.RData")
}
```

**Random Forests:** Typically you may want build to smaller trees

The classe variable is actually a categorical variable and therefore a classification method performs better. One could use a single tree, but Random Forest have proven to be the most accurate classification algorithm, mainly for the reduction of variability while averaging different random trees. The out-of-bag (oob) error rate is important in this model:

In random forests, there is no need for cross-validation or a separate test set to get an unbiased estimate of the test set error. It is estimated internally, during the run, as follows: each tree is constructed using a different bootstrap sample from the original data. About one-third of the cases are left out of the bootstrap sample and not used in the construction of the kth tree. Put each case left out in the construction of the kth tree down the kth tree to get a classification. In this way, a test set classification is obtained for each case in about one-third of the trees. At the end of the run, take j to be the class that got most of the votes every time case n was oob. The proportion of times that j is not equal to the true class of n averaged over all cases is the oob error estimate. This has proven to be unbiased in many tests.

**Method 1:** Standard Random Forest model Run the model: [make sure parallel is running]

```
registerDoParallel()
Trfor1<- system.time(rf1 <- randomForest(classe ~ .,
                                         data = smalltraining[,colIndex],
                                         importance=TRUE))
```

Check Predictions:

```
rf1 # OOB estimate of error rate: 0.77%
```

```
##
## Call:
## randomForest(formula = classe ~ ., data = smalltraining[, colIndex], importance = TRUE)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 6
##
##           OOB estimate of error rate: 0.77%
## Confusion matrix:
##           A      B      C      D      E class.error
## A 4460      4      0      0      0 0.0008961
## B  18 3009     11      0      0 0.0095458
## C   0  22 2710      6      0 0.0102264
## D   0   2  42 2527      2 0.0178780
## E   0   0   3  11 2872 0.0048510
```

```
rf1predictions1 <- predict(rf1, crossvalidation)
confusionMatrix(rf1predictions1,crossvalidation$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
```



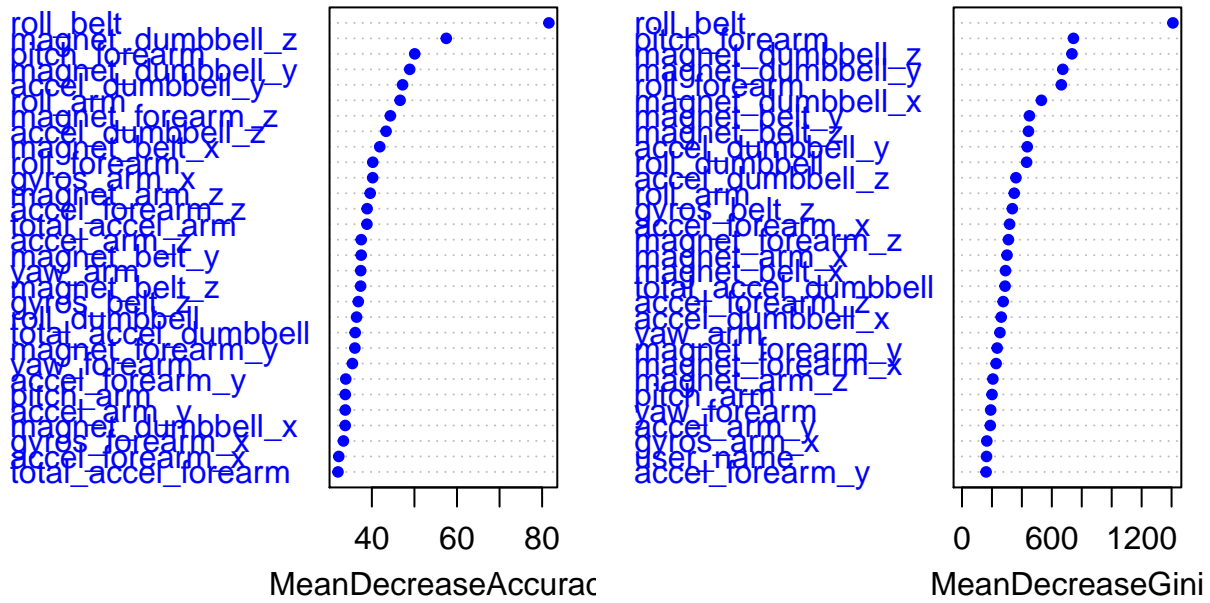
```
## Prediction      A      B      C      D      E
##           A 1114      2      0      0      0
##           B      2  754      6      0      0
##           C      0      3  678     11      0
##           D      0      0      0  632      2
##           E      0      0      0      0  719
##
## Overall Statistics
##
##           Accuracy : 0.993
##           95% CI : (0.99, 0.996)
##           No Information Rate : 0.284
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.992
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.998    0.993    0.991    0.983    0.997
## Specificity      0.999    0.997    0.996    0.999    1.000
## Pos Pred Value   0.998    0.990    0.980    0.997    1.000
## Neg Pred Value    0.999    0.998    0.998    0.997    0.999
## Prevalence       0.284    0.193    0.174    0.164    0.184
## Detection Rate    0.284    0.192    0.173    0.161    0.183
## Detection Prevalence 0.284    0.194    0.176    0.162    0.183
## Balanced Accuracy 0.999    0.995    0.993    0.991    0.999
```

```
rfor1<- confusionMatrix(rf1predictions1,crossvalidation$classe)
```

Assesment: What are the most influential trees? \*Exclusive to random forest

```
varImpPlot(rf1,pch=20,col="blue")
```


rf1



Plot: Choosing the right number of trees

```
plot(rf1, log="y")
legend("topright", colnames(rf1$err.rate), col=1:4, cex=0.8, fill=1:6)
```

rf1



**Method 2:** We now build 6 random forests with 150 trees each. We make use of parallel processing to build this model. Note: error with graphing tree

Set up and train model

```
t <- smalltraining[, colIndex]
x<- t[, -38]
y <- smalltraining$classe
Trfor2<- system.time(rf2 <- foreach(ntree=rep(150, 6),
                                   .combine=randomForest::combine,
                                   .packages='randomForest')
                    %dopar% {
                      randomForest(x, y, ntree=ntree)
                    })
```

Check Prediction

```
rf2 # OOB rate of 0% and used 900 trees
```

```
##
## Call:
## randomForest(x = x, y = y, ntree = ntree)
##           Type of random forest: classification
##           Number of trees: 900
## No. of variables tried at each split: 6
```

```
#we need to remove the missing values for this setup in
NonNAIndex <- which(colSums(is.na(crossvalidation)) > 0)
cross <- crossvalidation[ ,-NonNAIndex]

# cross is corssvalidation with missing variables missing
rf2predictions <- predict(rf2, cross)
confusionMatrix(rf2predictions,cross$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1116    0    0    0    0
##           B    0  759    0    0    0
##           C    0    0  684    0    0
##           D    0    0    0  643    0
##           E    0    0    0    0  721
##
## Overall Statistics
##
##           Accuracy : 1
##           95% CI : (0.999, 1)
## No Information Rate : 0.284
## P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 1
```

```
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      1.000    1.000    1.000    1.000    1.000
## Specificity      1.000    1.000    1.000    1.000    1.000
## Pos Pred Value   1.000    1.000    1.000    1.000    1.000
## Neg Pred Value   1.000    1.000    1.000    1.000    1.000
## Prevalence       0.284    0.193    0.174    0.164    0.184
## Detection Rate   0.284    0.193    0.174    0.164    0.184
## Detection Prevalence 0.284    0.193    0.174    0.164    0.184
## Balanced Accuracy 1.000    1.000    1.000    1.000    1.000
```

```
rfor2<- confusionMatrix(rf2predictions,crossvalidation$classe)
# 100% accurate?

#testing respect to original test set
rf2predictions2 <- predict(rf2, RemoveNA)
confusionMatrix(rf2predictions2,RemoveNA$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A    B    C    D    E
##           A 4464    0    0    0    0
##           B    0 3038    0    0    0
##           C    0    0 2738    0    0
##           D    0    0    0 2573    0
##           E    0    0    0    0 2886
```

```
## Overall Statistics
##
##           Accuracy : 1
##           95% CI : (1, 1)
##           No Information Rate : 0.284
##           P-Value [Acc > NIR] : <2e-16
```

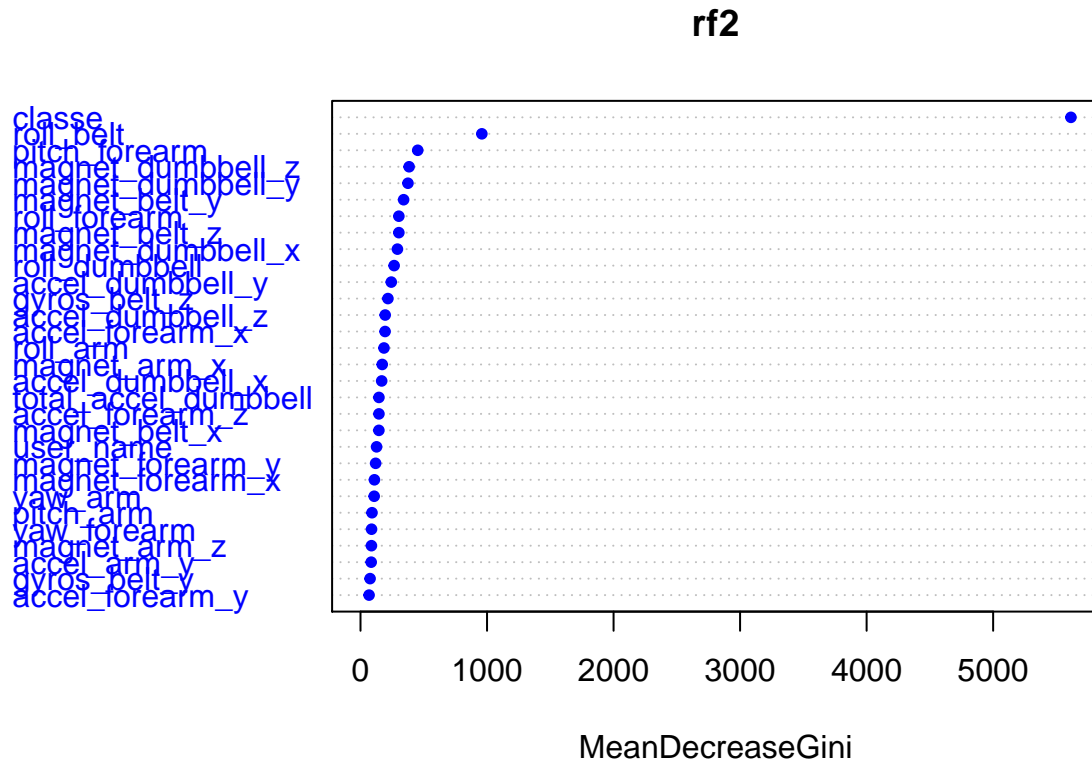
```
##           Kappa : 1
## McNemar's Test P-Value : NA
```

```
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      1.000    1.000    1.000    1.000    1.000
## Specificity      1.000    1.000    1.000    1.000    1.000
## Pos Pred Value   1.000    1.000    1.000    1.000    1.000
## Neg Pred Value   1.000    1.000    1.000    1.000    1.000
## Prevalence       0.284    0.194    0.174    0.164    0.184
## Detection Rate   0.284    0.194    0.174    0.164    0.184
## Detection Prevalence 0.284    0.194    0.174    0.164    0.184
## Balanced Accuracy 1.000    1.000    1.000    1.000    1.000
```

```
#100% accurate?
```

Assesment: What are the most influential trees?

```
varImpPlot(rf2,pch=20,col="blue")
```



**SVM RADIAL:** Support vector Machine is used for both classification and logistic regression. The radial kernal uses shortest distance of Euclidean distance. The kfoldes separate the sample in two and then the model trains each section to predict the other; this information is then used to create the final model. Increased folds may increase validity but for each increased fold there is less data to predict the model. So be careful

- I customized `train` control function to perform k-fold cross validation of 2.

Set up and run the model:

```
tC <- trainControl(method = "cv", number = 2) # note 'cv' creates folds and 2 is the size
TSVMRad<- system.time(SVMRadial1 <- train(classe ~ .,
                                          method = "svmRadial",
                                          trControl = tC,
                                          data = smalltraining[, colIndex]))
```

Check predictions:

```
SVMRadpredictions1 <- predict(SVMRadial1, crossvalidation)
confusionMatrix(SVMRadpredictions1,crossvalidation$classe)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction    A    B    C    D    E
##           A 1107   72    1    3    5
##           B    5  651   38    4   17
##           C    4   28  638   73   36
##           D    0    2    2  563   19
##           E    0    6    5    0  644
##
## Overall Statistics
##
##           Accuracy : 0.918
##           95% CI : (0.909, 0.927)
##           No Information Rate : 0.284
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.897
##           McNemar's Test P-Value : <2e-16
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity       0.992   0.858   0.933   0.876   0.893
## Specificity       0.971   0.980   0.956   0.993   0.997
## Pos Pred Value    0.932   0.910   0.819   0.961   0.983
## Neg Pred Value    0.997   0.966   0.985   0.976   0.976
## Prevalence        0.284   0.193   0.174   0.164   0.184
## Detection Rate    0.282   0.166   0.163   0.144   0.164
## Detection Prevalence 0.303   0.182   0.199   0.149   0.167
## Balanced Accuracy 0.982   0.919   0.945   0.934   0.945
```

```
SVMRad <- confusionMatrix(SVMRadpredictions1,crossvalidation$classe)
```

**SVM RADIAL COST:** Similar to above but it now implements a penalty to reduce possibility of overfitting

Setup and the run the model:

```
# model creation and test
tC <- trainControl(method = "cv", number = 2)
TSVMRadCost<- system.time(SVMRadialCost1 <- train(classe ~ .,
                                                    method = "svmRadialCost",
                                                    trControl = tC,
                                                    data = smalltraining[, colIndex]))
```

Check predictions:

```
SVMRadCostpredictions1 <- predict(SVMRadialCost1, crossvalidation)
confusionMatrix(SVMRadCostpredictions1,crossvalidation$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
```

```
## Prediction      A      B      C      D      E
##           A 1106    72     1     3     5
##           B     5   650    38     4    17
##           C     4    28   638    73    36
##           D     1     2     2   563    19
##           E     0     7     5     0   644
##
## Overall Statistics
##
##           Accuracy : 0.918
##           95% CI : (0.909, 0.926)
##           No Information Rate : 0.284
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.896
##           McNemar's Test P-Value : <2e-16
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity           0.991    0.856    0.933    0.876    0.893
## Specificity           0.971    0.980    0.956    0.993    0.996
## Pos Pred Value        0.932    0.910    0.819    0.959    0.982
## Neg Pred Value        0.996    0.966    0.985    0.976    0.976
## Prevalence            0.284    0.193    0.174    0.164    0.184
## Detection Rate        0.282    0.166    0.163    0.144    0.164
## Detection Prevalence  0.303    0.182    0.199    0.150    0.167
## Balanced Accuracy     0.981    0.918    0.945    0.934    0.945
```

```
SVMRadCost<- confusionMatrix(SVMRadCostpredictions1,crossvalidation$classe)
```

**TREE BAG:** it builds an expansive bundle of classification trees

Setup model and train it:

```
tC <- trainControl(method = "cv", number = 2)
TTB<- system.time(treebag1 <- train(classe ~ .,
                                   method = "treebag",
                                   trControl = tC,
                                   data = smalltraining[, colIndex]))
```

Check predictions:

```
treepredictions1 <- predict(treebag1, crossvalidation)
confusionMatrix(treepredictions1,crossvalidation$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      A      B      C      D      E
##           A 1110     9      0      2      0
##           B     4   736     9      1      1
```

```
##           C      0      9 670    18      2
##           D      0      5      3 620      4
##           E      2      0      2      2 714
##
## Overall Statistics
##
##           Accuracy : 0.981
##           95% CI : (0.977, 0.985)
##           No Information Rate : 0.284
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.976
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.995    0.970    0.980    0.964    0.990
## Specificity      0.996    0.995    0.991    0.996    0.998
## Pos Pred Value   0.990    0.980    0.959    0.981    0.992
## Neg Pred Value   0.998    0.993    0.996    0.993    0.998
## Prevalence       0.284    0.193    0.174    0.164    0.184
## Detection Rate   0.283    0.188    0.171    0.158    0.182
## Detection Prevalence 0.286    0.191    0.178    0.161    0.184
## Balanced Accuracy 0.995    0.982    0.985    0.980    0.994
```

```
TB<- confusionMatrix(treepredictions1,crossvalidation$classe)
```

Classification Tree: The most simplification form fo the classification tree

```
# apply classification tree
TCT<- system.time(Classtree1 <- train(classe ~ .,
                                     method="rpart",
                                     data = smalltraining[, colIndex]))
```

```
Classtreepredictions1 <- predict(Classtree1, crossvalidation)
confusionMatrix(Classtreepredictions1, crossvalidation$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  A    B    C    D    E
##           A 999 321 327 280 106
##           B  28 247  26 108  91
##           C  84 191 331 255 201
##           D   0   0   0   0   0
##           E   5   0   0   0 323
##
## Overall Statistics
##
##           Accuracy : 0.484
##           95% CI : (0.469, 0.5)
##           No Information Rate : 0.284
```



```
##      P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.326
##  McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: A Class: B Class: C Class: D Class: E
## Sensitivity          0.895   0.325   0.4839   0.000   0.4480
## Specificity          0.632   0.920   0.7743   1.000   0.9984
## Pos Pred Value       0.491   0.494   0.3117   NaN    0.9848
## Neg Pred Value       0.938   0.850   0.8766   0.836   0.8893
## Prevalence           0.284   0.193   0.1744   0.164   0.1838
## Detection Rate       0.255   0.063   0.0844   0.000   0.0823
## Detection Prevalence 0.518   0.127   0.2707   0.000   0.0836
## Balanced Accuracy     0.763   0.623   0.6291   0.500   0.7232
```

```
CT <- confusionMatrix(Classesreepredictions1, crossvalidation$classe)
```

Gradient Boosting (GBM)- is a machine learning technique for regression problems, which produces a prediction model in the form of an ensemble of weak prediction models, typically decision trees. It builds the model in a stage-wise fashion like other boosting methods do, and it generalizes them by allowing optimization of an arbitrary differentiable loss function. The gradient boosting method can also be used for classification problems by reducing them to regression with a suitable loss function.

Take a smaller sample to train model:

```
sampletrain <- smalltraining[sample(nrow(smalltraining), 3000), ]
inTrain <- createDataPartition(y=sampletrain$classe, p=0.7, list=FALSE)
tinytraining <- sampletrain[inTrain, ]
tinycrossvalidation <- sampletrain[-inTrain, ]
```

Set up the grid and run the model:

```
gbmGrid <- expand.grid(interaction.depth = 5, # the num of interactions between features
                      n.trees = 150, # the total number of trees or iterations
                      shrinkage = 0.1) # the learning rate of step-size function
TGBM<- system.time(GBM1 <- train(classe ~ .,
                                method="gbm",
                                data=tinytraining[,colIndex],
                                tuneGrid = gbmGrid,
                                verbose = FALSE)) #verbose doesn't show output iterations
```

Check predictions:

```
# test tiny crossvalidation
GBM1predictions<-predict(GBM1,tinycrossvalidation)
confusionMatrix(GBM1predictions,tinycrossvalidation$classe)
```

```
## Confusion Matrix and Statistics
##
```

```
##           Reference
## Prediction   A    B    C    D    E
##           A 247  10    0    0    1
##           B   1 160    9    0    0
##           C   0   7 146    9    2
##           D   0   1   5 133    4
##           E   0   0   1   1 160
##
## Overall Statistics
##
##           Accuracy : 0.943
##           95% CI : (0.926, 0.957)
##           No Information Rate : 0.276
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.928
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity           0.996   0.899   0.907   0.930   0.958
## Specificity           0.983   0.986   0.976   0.987   0.997
## Pos Pred Value        0.957   0.941   0.890   0.930   0.988
## Neg Pred Value        0.998   0.975   0.980   0.987   0.990
## Prevalence            0.276   0.198   0.179   0.159   0.186
## Detection Rate        0.275   0.178   0.163   0.148   0.178
## Detection Prevalence  0.288   0.190   0.183   0.159   0.181
## Balanced Accuracy     0.990   0.942   0.941   0.958   0.978
```

```
# test cross validation
GBM1predictions2<-predict(GBM1,crossvalidation)
confusionMatrix(GBM1predictions2,crossvalidation$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1094   37    1    1    5
##           B   12  671   44    8   17
##           C    5   42  620   37   10
##           D    4    5   18  592   16
##           E    1    4    1    5  673
##
## Overall Statistics
##
##           Accuracy : 0.93
##           95% CI : (0.922, 0.938)
##           No Information Rate : 0.284
##           P-Value [Acc > NIR] : < 2e-16
##
##           Kappa : 0.912
##           McNemar's Test P-Value : 5.33e-07
##
```

```
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.980   0.884   0.906   0.921   0.933
## Specificity      0.984   0.974   0.971   0.987   0.997
## Pos Pred Value   0.961   0.892   0.868   0.932   0.984
## Neg Pred Value   0.992   0.972   0.980   0.984   0.985
## Prevalence       0.284   0.193   0.174   0.164   0.184
## Detection Rate   0.279   0.171   0.158   0.151   0.172
## Detection Prevalence 0.290 0.192 0.182 0.162 0.174
## Balanced Accuracy 0.982   0.929   0.939   0.954   0.965
```

```
GBM<- confusionMatrix(GBM1predictions2,crossvalidation$classe)
```

## Comparing Models

Measuring Accuracy and Out of Sample Error:

```
# sum up all the methods
FinalAccuracy <- data.frame(rfor1$overall[1], rfor2$overall[1], SVMRad$overall[1],
                           SVMRadCost$overall[1], TB$overall[1], CT$overall[1],
                           GBM$overall[1])
colnames(FinalAccuracy) <- c("rfor1", "rfor2", "SVMRad", "SVMRadCost", "TB", "CT", "GBM")
rownames(FinalAccuracy) <- "Accuracy"
FinalAccuracy
```

```
##           rfor1 rfor2 SVMRad SVMRadCost      TB      CT      GBM
## Accuracy 0.9934      1 0.9184      0.9179 0.9814 0.4843 0.9304
```

```
# show the out-of-sample error
outOfSamErr <- 1-FinalAccuracy
rownames(outOfSamErr) <- "OSError"
outOfSamErr
```

```
##           rfor1 rfor2 SVMRad SVMRadCost      TB      CT      GBM
## OSError 0.006628      0 0.08157      0.08208 0.01861 0.5157 0.06959
```

Measuring Kappa- Goodness to Fit

```
FinalKappa <- data.frame(rfor1$overall[2], rfor2$overall[2], SVMRad$overall[2],
                        SVMRadCost$overall[2], TB$overall[2], CT$overall[2],
                        GBM$overall[2])
colnames(FinalKappa) <- c("rfor1", "rfor2", "SVMRad", "SVMRadCost", "TB", "CT", "GBM")
rownames(FinalKappa) <- "Kappa"
FinalKappa
```

```
##           rfor1 rfor2 SVMRad SVMRadCost      TB      CT      GBM
## Kappa 0.9916      1 0.8966      0.8959 0.9765 0.3262 0.9119
```

Measuring Size of each prediction model

```
FinalSize <- data.frame(format(object.size(rf1), units = "MB"),
                        format(object.size(rf2), units = "MB"),
                        format(object.size(SVMRadial1), units = "MB"),
                        format(object.size(SVMRadialCost1), units = "MB"),
                        format(object.size(treebag1), units = "MB"),
                        format(object.size(Classtree1), units = "MB"),
                        format(object.size(GBM1), units = "MB"))
colnames(FinalSize) <- c("rfor1", "rfor2", "SVMRad", "SVMRadCost", "TB", "CT", "GBM")
rownames(FinalSize) <- "Size"
FinalSize
```

```
##           rfor1  rfor2 SVMRad SVMRadCost      TB      CT      GBM
## Size 31.3 Mb 54.3 Mb 9.4 Mb      9.4 Mb 187.8 Mb 12.6 Mb 3.5 Mb
```

Comparing computation time:

```
FinalComp <- rbind(Trfor1, Trfor2, TSVMRad, TSVMRadCost, TTB, TCT, TGBM)
rownames(FinalComp) <- c("rfor1", "rfor2", "SVMRad", "SVMRadCost", "TB", "CT", "GBM")
FinalComp
```

```
##           user.self sys.self elapsed user.child sys.child
## rfor1           95.100    1.620   113.29      0.00    0.000
## rfor2            1.737    1.110   158.65     66.99   11.475
## SVMRad          32.607    5.600   521.33     70.22    2.798
## SVMRadCost      30.673    4.399   468.57    195.30   33.380
## TB              52.930    1.151   138.17     95.85   13.183
## CT               7.076    1.056   365.35     54.72    4.868
## GBM              5.978    0.220    89.47     94.46    3.838
```

Complete Model Comparison:

```
Group <- rbind(FinalKappa, outOfSamErr, FinalSize)
TGroup<- data.frame(t(Group)) # transform to matrix and transpose it
CompleteComparison<- data.frame(cbind(TGroup,FinalComp))
CompleteComparison <- mutate(CompleteComparison, usertime=user.self + user.child,
                             systime=sys.self + sys.child,
                             model = c("rfor1", "rfor2", "SVMRad", "SVMRadCost",
                                         "TB", "CT", "GBM"))
CompleteComparison<- CompleteComparison[, -c(4,5,7,8)]
CompleteComparison[, 1] <- round(as.numeric(as.character(CompleteComparison[, 1])), 3)
CompleteComparison[, 2] <- round(as.numeric(as.character(CompleteComparison[, 2])), 3)
CompleteComparison <- select(CompleteComparison,7,1:6)
CompleteComparison <- arrange(CompleteComparison, OSErrors)
CompleteComparison # without timestamp variables
```

```
##           model Kappa OSErrors      Size elapsed usertime systime
## 1      rfor2 1.000    0.000 54.3 Mb 158.65     68.73 12.585
## 2      rfor1 0.992    0.007 31.3 Mb 113.29     95.10  1.620
## 3         TB 0.976    0.019 187.8 Mb 138.17    148.78 14.334
## 4         GBM 0.912    0.070  3.5 Mb  89.47    100.44  4.058
## 5      SVMRad 0.897    0.082  9.4 Mb 521.33    102.82  8.398
## 6 SVMRadCost 0.896    0.082  9.4 Mb 468.57    225.97 37.779
## 7         CT 0.326    0.516 12.6 Mb 365.35     61.80  5.924
```

**Sub-comparison of excluding timestamp variables** Sub-Comparison to when we include timestamp variables [it was significant]

first

##	models	Kappa	OSError	Size	elapsed	usertime
## 1	rfor2	1.000	0.000	46.3 Mb	156.82	77.85
## 2	rfor1	0.998	0.001	26.9 Mb	146.02	94.88
## 3	TB	0.989	0.009	189.5 Mb	86.87	148.61
## 4	GBM	0.979	0.016	3.6 Mb	90.56	111.69
## 5	SVMRad	0.904	0.075	9.8 Mb	970.34	173.54
## 6	SVMRadCost	0.903	0.076	9.9 Mb	286.08	125.25
## 7	CT	0.326	0.516	13 Mb	48.24	111.21

Notice including timestamp variables decrease total size for almost all algorithms. Some models are impacted significantly computationally when including them while others enjoy one less variable. The reason is that it provides more possibilities to match and separate the data OR it makes it easy to reach the goal because of fewer variables. REGARDLESS it is worth to explore tradeoffs in more well-tweaked models.

**Conclusion** I've tested many machine learning models in this exercise. Normally, we just choose the most accurate algorithm and move on, but we need to consider the entire pipeline of the project. Several factors that we should care about is accuracy/out of sample error, fitted train model size, elapsed and system time. With that said the top three models are randomforest models, general boosting models, and treebag model. The treebag requires so much data to hold, so let's discard that. The GBM and randomforest are very both good candidates. Randomforest models have many additional features that shed a lot more of the internal processes, which can allow to build a more efficient model (less trees or remove the least interesting features). In contrast, GBM excels greatly in minimum size and training time while still maintaining accuracy.

With the comparison chart feature you can do short diagnostic on which model you want to implement. Note the logistic regression would use a similar procedure with a few differences.

**Bibliography** Velloso, E.; Bulling, A.; Gellersen, H.; Ugulino, W.; Fuks, H. Qualitative Activity Recognition of Weight Lifting Exercises. Proceedings of 4th International Conference in Cooperation with SIGCHI (Augmented Human '13) . Stuttgart, Germany: ACM SIGCHI, 2013.

**Additional Material** **Principal Component Analysis-** It can only handle numeric vectors

```
preProcompact <- preProcess(compacttraining[,-c(1,57)], method="pca")
preProcdescriptive <- preProcess(descriptivetraining[,-c(1,42)], method="pca")
```

Do NOT Use PCA to model fit for large datasets as it crashed for R or taken enormous amount of time to complete.

**Per person approach**

**Train classifier for training subset** Based on the findings from the previous section, we'll learn separate predictor for each user.

**Apply model to test subset and determine out of sample error** The out-of sample error seems to be well under control.

**Train classifier on full test set**